

Analysis Documentation

Student Performance Prediction

Prepared by Kavın Ganapathy

Student ID: 100008820

Date: 2026-06-07

Table of Contents

1. Introduction	3
2. Descriptive Analysis	4
3. Categorical Testing	6
4. Correlation & Relationships	7
5. Predictive Trends — Linear Regression	8
6. Synthesis & Conclusions	9

1. Introduction

1.1 Purpose of this document

This document presents a structured statistical analysis of the **Student Performance Prediction** dataset. It is written to be read top-to-bottom: each section poses a concrete question, states the method used to answer it, presents the evidence (a chart or table), and closes with an interpretation. The goal is a defensible, reproducible understanding of what drives student outcomes.

1.2 Problem statement

PROBLEM STATEMENT

Educational institutions need to understand which study-related factors are associated with student exam performance, whether outcomes differ across student groups, and whether performance can be predicted from measurable behaviours. This study interrogates the Student Performance Prediction dataset (10,000 students, 8 variables) to answer those questions.

1.3 Objectives

- Describe the central tendency, spread and shape of the key variables.
- Test whether exam performance differs significantly across categorical groups.
- Measure which factors are most strongly correlated with performance.
- Build and evaluate a predictive linear-regression model.
- Synthesise the findings into an overall understanding of the dataset.

1.4 Methodology & tools

Analyses were produced with the Smart Analysis Reporter toolkit (Python, pandas, SciPy, statsmodels). Descriptive statistics, frequency analysis, correlation, independent-samples t-tests and ordinary-least-squares regression were applied, each with a significance level of $\alpha = 0.05$.

2. Descriptive Analysis

PROBLEM STATEMENT

What are the typical values, variability and distribution shape of the study habits and performance measures in the dataset?

2.1 Method

Summary statistics (mean, median, mode, variance, standard deviation, quartiles, skewness and kurtosis) were computed for every metric variable.

2.2 Results

	study_hours	attendance	sleep_hours	internet_usage	assignments_completed	project_score	exam_score
N	10000	10000	10000	10000	10000	10000	10000
Mean	5.9896	69.8846	6.4985	6.0626	9.9884	64.911	86.7042
Median	6.0	70.0	6.5	6.0	10.0	65.0	92.12
Mode	10	56	5	5	5	57	100.0
Variance	10.0083	310.3421	2.9219	9.8481	36.4109	306.3557	226.7549
Std	3.1636	17.6165	1.7094	3.1382	6.0341	17.503	15.0584
Min	1	40	4	1	0	35	26.67
Max	11	100	9	11	20	95	100.0
Range	10	60	5	10	20	60	73.33
Q1	3.0	55.0	5.0	3.0	5.0	50.0	76.7275
Q3	9.0	85.0	8.0	9.0	15.0	80.0	100.0
IQR	6.0	30.0	3.0	6.0	10.0	30.0	23.2725
Skewness	-0.0072	0.0178	0.0051	-0.0135	0.0178	0.0004	-1.0086
Kurtosis	-1.2218	-1.2075	-1.2714	-1.2008	-1.1959	-1.1856	0.1475
Missing	0	0	0	0	0	0	0




Figure 2.1 — Distribution of the target variable, exam_score.

2.3 Interpretation

exam_score averages 86.70 (median 92.12, SD 15.06), ranging from 26.67 to 100.00. The distribution is left-skewed (skewness -1.01).

3. Categorical Testing

PROBLEM STATEMENT

Do students differ in exam performance across categorical groups — for example, between placement outcomes?

3.1 Group composition




Figure 3.1 — Distribution of placement_status.

placement_status has 2 categories; the most common is "Placed" (83.56% of records).

3.2 Hypothesis test (independent-samples t-test)

H_0 : the mean exam_score is equal across placement_status groups. H_a : the means differ.

placement_status	Mean	SD	N
Placed	91.997	9.377	8356.000
Not Placed	59.805	8.147	1644.000

3.3 Interpretation

Reject H_0 — exam_score differs significantly between "Placed" and "Not Placed" ($t = 142.69$, $p < 0.001$, Cohen's $d = 3.50$, large effect).

4. Correlation & Relationships

PROBLEM STATEMENT

Which variables are most strongly associated with exam performance, and could therefore serve as predictors?

4.1 Method

Pearson correlation coefficients were computed for all metric variable pairs (range -1 to +1).

4.2 Results

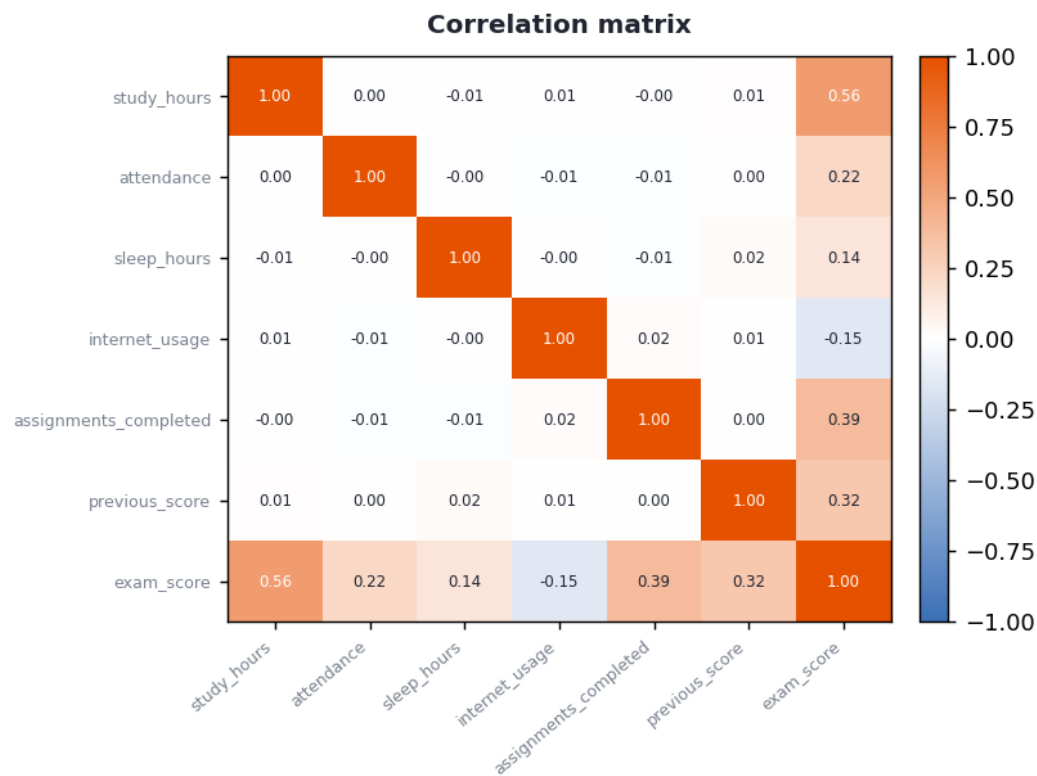


Figure 4.1 — Correlation matrix of the metric variables.

4.3 Interpretation

The strongest linear relationship is between study_hours and exam_score ($r = 0.56$), a moderate positive association. Variables with high correlation may carry overlapping information.

5. Predictive Trends — Linear Regression

PROBLEM STATEMENT

Can exam performance be predicted from its strongest single factor, and how much of the variation does that factor explain?

5.1 Method

A simple ordinary-least-squares linear regression was fitted, modelling exam_score as a function of study_hours.

5.2 Results




Figure 5.1 — Regression of exam_score on study_hours.

5.3 Interpretation

A simple linear model predicts exam_score from study_hours as $Y = 70.667 + 2.678 \times \text{study_hours}$. It explains 31.6% of the variance ($R^2 = 0.316$). The slope is statistically significant ($p < 0.001$), so each one-unit rise in study_hours is associated with a change of 2.678 in exam_score.

As a worked example, a student one standard deviation above the mean study_hours (9.2) is predicted to score **95.2** on exam_score.

6. Synthesis & Conclusions

6.1 Bringing the analysis together

Each section answered one question, and together they form a coherent picture. The descriptive profile set expectations; categorical testing showed where groups genuinely differ; correlation identified the factors that move with performance; and the regression turned the strongest of those into a predictive rule.

6.2 Key findings

- The strongest linear relationship is between **study_hours** and **exam_score** ($r = 0.56$).
- A simple model of **exam_score** on **study_hours** explains 31.6% of its variance ($R^2 = 0.32$).
- There is a statistically significant difference in **exam_score** between **placement_status** groups ($p < 0.001$).
- **exam_score** is not normally distributed (Shapiro-Wilk $W = 0.8376$).

6.3 Limitations & next steps

- A single-predictor model is a baseline — a multiple-regression model using several factors together would likely explain more variance.
- Correlation does not imply causation; a controlled study is needed to confirm drivers.
- Findings describe this sample; external validation on new cohorts is recommended.